

THE UMA-SOM PLATFORM FOR GENERAL LEARNING

– A USER’S MANUAL –

DAN P. GURALNIK[†] AND SIQI HUANG

ABSTRACT. This technical report supplements a document by the same name describing the research and development effort by our group on a partially distributed, GPU-accelerated implementation of a modality-agnostic architecture for life-long general learning, based on the Universal Memory Architecture (UMA) memory model previously introduced by Guralnik and Koditschek. The design by Huang and Guralnik emphasizes modularity, distributability and flexibility of application, by separating the computational components of the architecture (implemented in CUDA) from the RESTful communication framework (C++) and the simulation state update and environment (Python).

This document is envisioned both as a user’s manual and as a technical description of the architecture in support of further development efforts. Finally, it also contains a running description of existing applications and experimental results.

1. PREPARING, INSPECTING AND APPLYING THE UMA-SOM PACKAGE

This section focuses on preparing the UMA-SOM system for work without paying much attention to system design. A minimal overview of the design at the highest level is provided to help the user follow the installation process and understand proper use. Knowledge of the parent document [1]—specifically Section 3.1 therein—is assumed for the purpose of reading Section 1.4, discussing the writing of simulation scripts.

1.1. Building the UMA Package. The UMA project consists of the following parts:

- (1) UMA Simulation Engine (Python),
- (2) UMA Client (Python),
- (3) UMA Core Service (C++).

The UMA Core Service is written in C++ and needs to be built by Visual Studio (on Windows) or by `gcc` (on Linux). UMA Core Service supports either a Windows or Linux build, and has been tested on Win7, Win10 and Ubuntu 16.04 LTS. UMA Core Service has three additional library dependencies, they are: `gtest` (Google Test), `cuda` (currently NVIDIA CUDA 8.0) and `cpprestsdk` (Microsoft Casablanca RESTful Interface). The `cuda` and `cpprestsdk` libraries require installation prior to building the project (e.g., using `apt-get` on Linux, or `vcpkg` on Windows); the `gtest` library is built in tandem with the UMA project, using a provided batch/shell script, depending on the OS. We also make use of the `Cmake` utility to streamline the build process for the user.

Detailed instructions for building are given in the project repository `readme.txt` file.

[†] UNIVERSITY OF PENNSYLVANIA, ELECTRICAL & SYSTEMS ENGINEERING DEPARTMENT, PHILADELPHIA, PA.

1.2. **Directory Tree.** The UMA folder has several folders and files:

- **bin/**

A successful build will result in the following files being written to the **bin** folder: the project executable file **UMA(.exe)**, and the shared (dynamic link) libraries **UMACore.so(dll)** and **UMARest.so(dll)**.

- **ini/**

The **ini/** folders under both **bin/** and **bin/test/** contains all the parameters necessary to launch the UMA program. Please refer to "UMA C++ Module/C++ ini files" for more details.

- **test/**

This folder contains other executable files used for C++ unit testing.

- **log/**

This folder will appear once UMA is run for the first time. Separate execution logs for each component of the project will be written to this folder at run time.

- **dependency/**

The source code in this library is built by the shell script **build.dependency.bat(sh)**. The folder houses all third party libraries used in the project:

- Google Test (**gtest**), used by the C++ unit testing component.

All the executables generated under **bin/test** depend on this library.

- **pylib/**

Contains the Python libraries for the UMA front end:

- **UMA/**

Libraries defining the UMA architecture data structures—Signals, Experiments, Agents—and some supporting functions. These are directly included by simulation scripts in the **suites/** folder. Their scope does not extend beyond the Python side.

- **Client/**

In contrast, the **Client** sub-folder contains Python scripts running outside the scope of Python. These are responsible for connecting to the core services.

- **perf/**

Finally, the **perf** sub-folder contains libraries for performance testing.

- **src/**

Contains all the C++ source code. Each sub-folder houses a C++ module, except the **common** directory:

- **common/**

All the frequently used headers reside here.

- **utility**

Contains all utility functionalities in the entire project. The corresponding C++ module is **UMAUtil**.

- **kernel/**

Contains CUDA kernel code and other functionalities requiring a GPU to run. **UMAKernel** is the resulting C++ module.

- **RESTAPI/**

Contains the entry of the project (main function) and the files that construct the REST API callable from the Python client.

- **RESTful/**

Includes the basic REST functionality used in the project, wrapped around **cpprestsdk**. Its C++ module is **UMARest**.

- **core/**

Source code for the module **UMACore**, implementing BUAs, their associated data structures and functionalities.

- **kernelTest**

Unit test files for module **UMAKernel**. Its C++ module is **UMAKernelTest**.

- **utilityTest**

Unit test files for module **UMAUtil**. Its C++ module is **UMAUtilTest**.

- **CMakeLists.txt**

- **suite/**

Contains the available Python simulation scripts using the **UMACore** service:

- **sniffy/**

See [1], Section 4.1.1 for a detailed description of this environment.

- **mice/**

See [1], Section 4.1.2 for a detailed description of this environment.

- **test/**

Contains all functional and performance test scripts.

- `build_dependency.bat`

Must be run to build the code under `dependency/` under Windows.

- `build_dependency.sh`

Must be run to build the code under `dependency/` under Ubuntu.

- `build_vs.bat`

This batch file is used on Windows only, to build the Visual Studio solution `UMA.sln`.

- `build_kernel.sh`

This shell script is used on Linux only, to build *only* the `UMAKernel` module (on Windows, `UMAKernel` is built by Visual Studio together with the other modules).

- `readme.txt`

Detailed instructions for preparing and running the build.

1.3. Deployment. An UMA deployment includes two parts: the UMA Python front end and the UMA Core (C++/CUDA back end), see Figure 1.

UMA Python modules can run on any platform supporting Python 2.7 (preferably an up-to-date Anaconda distribution); the UMA Core, for now, only supports Windows and Linux. Refer to Section 1.1 for details on building the UMA core.

For small simulations and/or debugging purposes (for example, see [1], Section 4.2.1), we recommend to run UMA Core on Windows, and build with MSVC. UMA Python modules should be run on the same machine, and use `localhost` as their server. For larger scale deployment, it is preferable to build and deploy the UMA core on Linux, and launch UMA Python on a separate instance (Linux also preferred), residing in the same region, to maintain low network traffic latency. AWS would be a desired place to host the environment.




FIGURE 1. Separation of the UMA packages into a Python front end and a C++/CUDA back end.

The UMA Core application is currently only capable of running as a single instance: running multiple UMA instances and having them communicate with each other is not yet supported. However, this distributed functionality is easily achievable, due to the UMA REST functionality, by altering the `UMARest` module, for it to both handle and send requests to other UMA instances.

1.4. UMA simulation Scripts.

1.4.1. *Running a simulation script.* Running an UMA simulation requires a Python script to import the UMA-SOM engine `som2.py` from the `pylib/UMA/` project directory, as well as the client modules from `pylib/Client/`. We recommend adding these to the `$PYTHONPATH` environment variable. Alternatively, these modules may be copied into the working directory.

The `UMA(.exe)` executable needs to be run prior to executing the simulation script. The directory containing it must also contain any other items present in the `bin/` directory at the end of the build process (with the exception of the `bin/test/`—you may not be interested in testing—and the `bin/log/`—logs are formed at run time—directories).

Please make sure the log files are purged from time to time (see also initialization files settings, Section 2.2.4).

1.4.2. *Designing a new simulation script.* An UMA simulation operates within the framework of an Experiment Python object, whose role is to maintain the simulation state and to communicate with the BUAs through Agent objects. Section 3.1 of the technical report provides a detailed discussion of the update procedure, implemented by the `update_state` method of the `Experiment` class. Technical details from a more advanced developer’s perspective are discussed in Section 2.7. The bare essentials to be aware of are:

Experiment: An `Experiment` object must be constructed, and it is convenient to create a reference to the decision stack:

- `EX=Experiment()`
- `id_dec='decision'`

Registration: Each state variable—or *measurable*—must be registered before it could be addressed by any other component of the experiment:

- `<id_meas>=EX.register(<meas_name>)`

When registering a measurable, the user provides a unique ID tag (MID for short), to be used for referring to values of this measurable in the experiment state. If a tag is not provided, a *uuid* is issued and used by the system instead. Binary sensors—those corresponding to the control values provided by BUAs included—are issued MIDs in pairs, using:

- `<id_meas>,<cid_meas>=EX.register_sensor(<meas_name>)`

Sensors *must* be registered in this way (rather than as two separate MIDs).

Experiment State: The experiment state is maintained as a dictionary of deques¹, `EX._STATE`, indexed by the registered MIDs. The experiment state may be accessed as follows:

- `EX.state.all()`: provides access to the entire state dictionary;
- `EX.this_state(<mid>,<delta=0>)`: returns the state of measurable `<mid>`, `<delta>` cycles ago (default is the current state);
- `EX.set_state(<mid>,<val>)`: push `<val>` to the head of the deque for measurable `<mid>` (sets the current state).

Supplying Definitions: Each measurable requires a definition of its state update function as a function of the state dictionary. An appropriate function object must be constructed, and passed as an argument to the relevant constructor, as detailed below.

Constructing Measurables: Plain measurables are constructed using:

- `EX.construct_measurable(<mid>,<definition>=None,<init_value>=None,<depth>=1,<decdep>=False)`:
A measurable with MID `<mid>` is added to the experiment:

- (1) The provided MID *must* be registered before this function is called.
- (2) `<definition>` is a function object taking a single input, named `state`, envisioned as a function of the experiment state dictionary.

¹<https://docs.python.org/2/library/collections.html#collections.deque>

- (3) `<init_value>` is a list with which to initialize this measurable's entry in the experiment state dictionary.
- (4) `<depth>+1` is the desired bound on the length of the deque representing `<mid>` in the experiment state dictionary (default is 2, to store the current state and the preceding state).
- (5) `<decdep>` is a Boolean value indicating whether `<mid>` will participate in the decision sweep (see [1], Section 3.1.2).

Constructing Sensors: Following the same syntax as the constructor for measurables, the constructor for sensors is:

- `EX.construct_sensor(<mid>, <definition>=None, <init_value>=None, <depth>=1, <decdep>=False)`:

The only difference in functionality is that this constructor sets up *two* measurables: one with MID `<mid>`, and one with MID `<mid+'*>`, corresponding to the negated sensor.

Constructing Agents: BUAs, in their dual role of both decision-makers and sensors, require a separate constructor:

- `BUA=EX.construct_agent(<mid>, <signal>, <definition>, <params>)`

The input entries for the constructor are:

- (1) `<mid>` is the same MID used to register the *sensor* corresponding to the BUA being constructed;
- (2) `<signal>` is a non-negative real-valued measurable (whose MID should be registered when the constructor is called), coinciding with the BUA's *motivational signal* (see [1], opening of Section 2.1.2);
- (3) `<definition>` is the BUA's sensor update function;
- (4) `<params>` are the BUAs initialization parameters: currently, a pair (Q, AT) , where $Q \in (\frac{1}{2}, 1)$ is the discount parameter and $AT = \text{True/False}$ depending on whether the BUA will be auto-targeting or not (see [1], Section 2.1.2, stage (6) of BUA decision cycle).

Assigning Sensors to Agents: Once an Agent class object BUA has been constructed, it is possible to assign sensors to it:

- `BUA.add_sensor(<mid>)`

Assigns the sensor `<mid>` to both snapshot structures of Agent BUA. Scripted sensors need to be assigned prior to initialization (see below).

After BUA has been initialized, a different syntax is available:

- `BUA.add_sensor(<mid>, <token>=None)`

The sensor is assigned to the snapshot specified by `<token>` (allowed values 'plus', 'minus' or unspecified, which defaults to the active snapshot). This is not intended for use in simulation scripts, and will interfere with proper function if used after BUA has been initialized (see below).

Initialization: Every BUA needs to be set up in the UMA core before it can be used. This is done using the instruction:

- `BUA.init()`

No scripted sensors may be added to BUA after this instruction is given.

State Update and BUA reporting: Once all BUAs have been initialized, one is free to advance the simulation clock. This is done by a call of the form:

- `<ex_report>, <agent_reports>= EX.update_state(<instruction>='decide')`

This is the main tool for gathering data and interacting with the simulation engine, where:

- (1) `<ex_report>` is a dictionary reporting any observations of the experiment update process;
- (2) `<agent_reports>` is a dictionary of reports produced by the individual BUAs during deliberation;
- (3) `<instruction>` is either absent, or is required to be a list of MIDs containing, for each scripted BUA, either its MID or the MID of its negation (but not both); when present, the instruction will be copied over the decisions produced by the scripted agents during the decision sweep (so as not to prevent any further deliberation such as—but not limited to—hard-wired arbitration of conflicting actions).

Tweaking the Reports: While most of the internal state of any BUA could be reported, any reporting puts an additional and unnecessary load on the communication framework. For this reason, constructing an adequate report for the purpose of any given simulation must be done by the user, by adding the appropriate instructions to the method `EX.update_state` of the module `UMA Python module, som2.py`.

The user should feel encouraged to stick to the established format, in the form of the dictionaries (1-2) above, while avoiding significant alterations to the core code of the `EX.update_state` method.

1.4.3. *Altering the UMA Python components.* Due to the global nature of the experiment state update, at this point, *the code in `EX.update_state` is extremely fragile*. Any alterations to the code beyond what is necessary for collecting the required data about the experiment requires a thorough understanding—not just of the algorithms—but also of the code base, at least as discussed in the next section.

2. DESIGN AND FUNCTION OF THE UMA-SOM PACKAGE

2.1. The UMA Architecture. Recalling Figure 1, we review the top-level components of the UMA architecture:

Python Simulation Engine, or “Front End” (FE): The Python simulation engine consists of the UMA Python modules (currently `som2.py`), and of the user’s executable simulation script. They reside, respectively, under `pylib/UMA/` and—for those simulation scripts we had developed—under `suite/`. Details on how to use the simulation engine are provided in Section 1.4 above, while the function of the UMA Python modules is reviewed in detail in [1], Section 3.

UMA Python Client: The Python client connects the simulation engine to the UMA core service. The source files reside under `pylib/client/`. The files are separated by their functionalities, for example: to request service from an Agent object, use `Service_Agent.py`; if a snapshot object is required to provide service, call `Service_Snapshot.py`. Further details are discussed in Section 2.7.1.

UMA Core Service, or “Back End” (BE): Written in C++, the UMA Core service supports CRUD (create/read/update/delete) of UMA objects (see Section 2.2); implementing BUA functionalities² (snapshot maintenance, inference and control); customized logging; and handling/throwing exceptions. The core service runs both on CPU and GPU. We cover the CPU/GPU data exchange in Section 2.3. The FE/BE communication protocol is discussed in Section 2.4.

2.2. UMA C++ Core.

2.2.1. C++ Modules. The UMA core services consist of seven C++ modules, whose dependency structure is summarized in Figure 2:

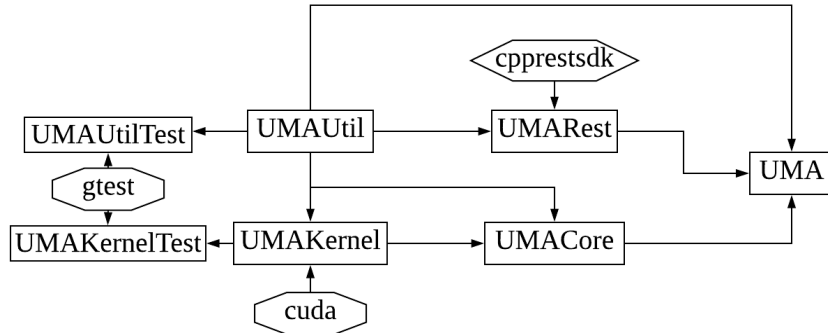


FIGURE 2. Package dependencies within the C++ UMA core.

(1) `UMAUtil.a(.lib)`

`UMAUtil` is a static library containing basic utility functions, the logging system, and exceptions consumed across the entire project. Every UMA module includes it as a base library. We recommend implementing any additional common functionalities under `UMAUtil`, if they are to be used often. The module itself only relies on the C++ library `std`, with no other dependencies. Hence, this module must be built first.

²See Section 2 of [1] for details.

(2) `UMAUtilTest(.exe)`

`UMAUtilTest` is an executable containing the unit test cases for the `UMAUtil` module. It depends on the `UMAUtil` and `gtest` modules. New test cases should be added here whenever a new functionality is implemented in `UMAUtil`.

(3) `UMAKernel.a(.lib)`

`UMAKernel` is a static library which contains all CUDA code (running on a Nvidia GPU device). The module consists of pure c++ files (`.cpp`), cuda files (`.cu`), and cuda headers (`.cuh`). All are compiled under `nvcc`.

All CUDA related syntax throughout the project is included here, including CUDA headers. The purpose of keeping this convention is to ensure `UMAKernel` is the only CUDA-dependent module.

Note that, despite calls to `UMAUtil` functionalities in `UMAKernel`, the latter may still be compiled separately because it is a static library.

(4) `UMAKernelTest(.exe)`

`UMAKernelTest` is an executable containing unit test cases for the `UMAKernel` module. It depends on `UMAKernel`, `gtest` and `UMAUtil` (see preceding item).

(5) `UMACore.so(.dll)`

`UMACore` is a shared (dynamic linking) library, that contains all UMA simulation objects. The objects in `UMACore` are consistent with the UMA simulation entities, so that UMA Python Simulation can operate directly on those objects by using REST API. The module uses `UMAKernel` functionalities but employs no CUDA syntax. As a result, only the host compiler is required for this module, with no dependency on CUDA.

The object classes defined in this module are key to the entire project, and we review them here:

World: a global unique object, that holds agents, it is the top level structure in `UMACore`. It is accessible from the Python REST API.

Agent: implements the functionalities of a BUA. A REST API is provided which operates on agents directly. Currently, an Agent maintains two snapshots.

Snapshot: implements a snapshot entity. A REST API is provided which operates on Snapshot objects directly. A Snapshot refers to its associated collections of Sensor, SensorPair, AttrSensor and AttrSensorPair objects (see below). It also refers to a DataManager object. Snapshot objects handle an overwhelming majority of service requests, but they do not handle the CPU/GPU data flow directly.

DataManager: this object type is essential for the data flow between GPU, CPU and all UMA objects. It is responsible for creating data arrays on both host and device, resizing those arrays, extracting them for computations, and cleaning them up. All low-level related data flow should use this object. For more info about the data flow in UMA, see Section 2.3.

Sensor: represents a sensor (Boolean query) entity assigned to the BUA (the parent Agent object), and it is accessible directly from the client side. It holds a ‘positive’ and a ‘negative’ **AttrSensor** objects (corresponding to the actual query and its negation), and refers to variables stored in the BUA’s **DataManager** object. During resizing operations, it is also responsible for storing some state variables.

SensorPair: Responsible for recording the information pertinent to learning and reasoning, such as the snapshot weights (see, e.g. [1] Figure 1) and learning thresholds. It controls 4 **AttrSensorPair** objects, and has pointers to the two **Sensor** objects making up the pair. During resizing, it is also responsible for storing the learning threshold value for the pair.

AttrSensor: This objects stores attributes of the sensor from the vantage point of the parent Agent object. It stores the current observation (state of the sensor), and its current and preceding weights. It is accessible from REST API on client side.

AttrSensorPair: For every pair of **AttrSensor** objects, there is a **AttrSensorPair** object to manage them, which is controlled by the appropriate **SensorPair** object. An **AttrSensorPair** object stores the weight of an edge in the poc graph and the orientation of the edge (if any). It also refers to the **SensorPair** objects making the pair.

Finally, there is a dedicated name space for all operations on the above objects. See Section 2.5 for further details.

(6) **UMARest.so(.dll)**

UMARest is another shared (dynamic linking) library, that only deals with REST functionality. It functions as a wrapper for **cpprestsdk** (Casablanca), to facilitate the construction of REST APIs without requiring knowledge of the 3rd party library. **UMARest** has three important object classes:

UMARestListener: Similarly to a server, its function is to catch input REST calls, and pass them to corresponding handlers.

UMARestHandler: This is the base class for all customized handlers. Upon receipt of a request from the listener, its task is to reply with a corresponding response.

UMARestRequest: An object of this class contains matching requests and response objects. It is tasked with holding all information about a request, and with converting input data into a format acceptable for **cpprestsdk**. It is also responsible for casting outgoing data into an appropriate type from the C++ **std** library, to facilitate communication with other UMA modeules.

(7) **UMA(.exe)**

This is the final executable. When it launches, it will listen on a local port, and uses **UMARest** to serve incoming requests. When a request comes in, **UMARest** will pass it to the appropriate handler and call the functionality in **UMACore**.

2.2.2. *C++ Exceptions.* The exception class resides in `UMAUtl` (which is included in all UMA modules), to ensure that any C++ exception raised by an execution of UMA can be handled in any of its modules. All Exceptions inherit from `std::runtime_error`, and are initialized with an error message, error level, and error type. Error messages are customizable. Implemented error levels are: `WARN`, `ERROR` and `FATAL`. `WARN` means the exception being raised did not interrupt the current workflow, and it is rarely used in current code; `ERROR` means the workflow has been compromised; `FATAL` indicates the operation interrupted the UMA process, causing its immediate termination. Error types classify the possible exceptions. When handling an exception, `UMARest` will return a status codes to the client, based on the error type.

2.2.3. *C++ logging system.* Similarly to the exception handling system, the logging system is implemented in `UMAUtl` to ensure all modules can utilize it. Each architectural component maintains its separate log, which can be set to a preferred level of detail: `ERROR`, `WARN`, `INFO`, `DEBUG`, `VERBOSE`. If the log level is `INFO`, then message logged as `DEBUG` and `VERBOSE` will not be recorded. Log files are created when loggers are created. The logger will read from `log.ini` to get the component's log level. All loggers are static global variables.

2.2.4. *C++ initialization files.* Initialization files are located under `ini/` and are essential for the program to run. There are, so far, four such files.

log.ini: Sets the log level information for each log component.

restmap.ini: Contains the REST map information. A simple way to disable a REST endpoint is comment(#) out the appropriate url.

server.ini: Contains the host and port information used to initialize the `UMARest` listener.

core.ini: Sets default parameters for UMA Core objects, such as the discount parameters and learning thresholds of Agent objects and the allocated memory expansion rate for Snapshot objects.

The content of a `.ini` file is organized by the architectural component relevant to the task. Each component's section of the file is headed by its name in square brackets. For each component, keys and key value pairs are used to store its settings:

log.ini: Component names will match the list of Logger object names; 'level' is the only accepted key, which maps to the log level.

restmap.ini: Component names are the available REST handlers. Each component lists the URLs to be used for REST mapping.

server.ini: Has the server port and host message, under component 'Server'.

core.ini: Components are the names of UMA object classes. The key value pairs set default values that to be used by these objects during initialization.

2.3. UMA Data Flow. The data flow in the UMA architecture is based on its separation into components presented in Figure 1.

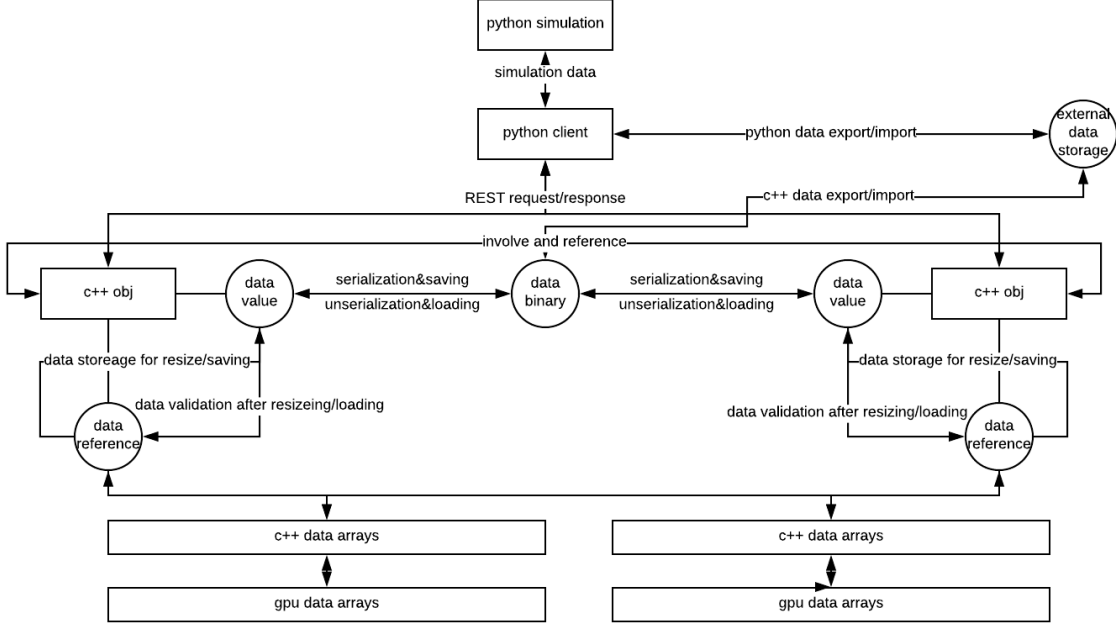


FIGURE 3. Data flow within the C++ UMA core.

On the Python side, the simulation engine will exchange data with the client APIs. The input from the simulation includes: current and preceding raw observation vectors for each BUA; current and preceding values of the motivational signal for each BUA; target specification vectors (for some BUAs) in some simulations; and, finally, BUA parameters. Output data are the decisions for each BUA; internal state vectors (target, prediction, current state,...) and results of batch-propagation.

As a result, the Python client and UMA C++ Objects data flow is rather complex, having many data types to deal with in the context of different REST calls. Please check the corresponding REST API code for details.

Data flow between UMA C++ Objects and Host/Device Data is another important layer.

- To leverage parallel computation capabilities, storage space must be allocated in continuous regions (arrays) both on CPU and GPU, for every data element;
- Pointers to all relevant data arrays must be maintained;
- All the above are complicated by the enrichment operation available to BUAs: extra memory needs to be allocated to accommodate dynamic extension of all data structures, with periodic resizing (expansion) whenever the allocated space budget is used up; resizing is achieved by allocating a larger continuous region to each array (both on CPU and on GPU), copying the contents of the old array into this region and freeing the old region, which forces an update of all the relevant pointers.

2.3.1. Other relevant functionalities. Work on a save/restore functionality of the BUA population in a given experiment had been started, but has not been completed at the time of writing of this document. The C++ implementation of any given BUA may be saved and recovered, but the global MID registration system of the Python simulation engine requires redesign to enable the stitching of a saved back-end UMA-SOM instance to a provided Python simulation script providing a compatible collection of initial scripted sensors (using shared MIDs).

2.4. The UMA REST Workflow. The UMA REST workflow represented in Figure 4 represents our implementation of **UMARest**. When a request is posted, **UMARest** will first locate the corresponding REST




FIGURE 4. The UMA REST workflow.

handler. Then it will scan the input fields for data. Finally, it will call the appropriate module in **UMACore** to execute the appropriate back end operation. If everything is OK, **UMARest** will return a 200+ response to the client, along with output data. Otherwise, **UMARest** will return a 400+/500+ error flag with error messages.

2.5. UMA REST Map. The UMA REST map has seven handlers, all read from the **restmap.ini** file. The handlers are:

WorldHandler: deals with agent management

- /UMA/world: return the current world's agent info

AgentHandler: manages Agent CRUD operations

- /UMA/object/agent: CRUD Agent

SnapshotHandler: manages Snapshot CRUD and maintenance

- /UMA/object/snapshot: manages CRUD in Snapshot
- /UMA/object/snapshot/init: initializes a snapshot using all current sensors as initial sensors
- /UMA/snpashot/delay: adds the specified delayed conjunctions to the snapshot
- /UMA/snapshot/amper: adds the specified conjunctions to the snapshot
- /UMA/snapshot/pruning: removes the specified delayed conjunctions from the snapshot

SensorHandler: manages Sensor/SensorPair CRUD

- /UMA/object/sensor: CRUD sensor
- /UMA/object/sensor_pair: CRUD sensor pair

AttrSensorHandler: manages AttrSensor/AttrSensorPair CRUD

- /UMA/object/attr_sensor: CRUD AttrSensor
- /UMA/object/attr_sensor_pair: CRUD AttrSensorPair

SimulationHandler: manages all simulation requests

- /UMA/simulation/decision: launch another decision cycle
- /UMA/simulation/up: perform upwards closure on a single input signal (deprecated)
- /UMA/simulation/ups: perform upwards closure on a batch of input signals
- /UMA/simulation/downs: perform downwards closure on a batch of input signals
- /UMA/simulation/propagation: perform propagation on a batch of input signals
- /UMA/simulation/npdirs: enforce the calculation of the implications matrix based on the current poc-graph matrix
- /UMA/simulation/blocks: perform single-linkage clustering at the specified scale on a batch of input signals
- /UMA/simulation/abduction: perform abduction on a batch of input signals
- /UMA/simulation/propagate_masks: enforce the propagation of the delay masks

DataHandler: manage all data arrays in core, for the host snapshot

- /UMA/data/observe: update the raw observation signal
- /UMA/data/current: update/read current state signal
- /UMA/data/prediction: read prediction signal
- /UMA/data/target: update/read target signal
- /UMA/data/data_size: read the current size info of all array variables
- /UMA/data/weights: report the weights matrix
- /UMA/data/thresholds: report the threshold matrix
- /UMA/data/dirs: report the implications matrix
- /UMA/data/negligible: report the indicator signal of the set of negligible sensors

2.6. UMA Test. The project implements three kinds of tests: unit, functional, and performance tests.

The unit testing utility is built using **gtest**. To expand the collection of unit tests, add the test cases to the C++ test source, and run the test executables for the module.

To run performance tests, run **pylib/perf/resource.py** as a separate process to have system matrix usage recorded as an experiment is being run.

Functional testing is based on REST API. Construct your functional test cases in yml form, and run with **test_wrapper.py** (which is found under **pylib/client/**). The functional test checks the status code, message and data fields arriving in response to the implemented REST calls.

2.7. Development Guide. This section includes comments on our design considerations and how they should affect further development by the user.

2.7.1. UMA Python Simulation Engine Development. The simulation engine uses the UMA Python Client to interact with the UMA Core. Simulation scripts may have varying settings, but will share some common libraries (such as `som2`). It is preferable for the simulation scripts not to interact with the UMA Python Client directly, so keeping track of the capabilities of the Agent class of `som2` and expanding them as necessary are recommended.

In the common libraries, to interact with UMA Python Client, an `UMA_service` object is (and must be) created in the global scope. The `UMA_service` object is a wrapper class that handles all REST calls to the the UMA core, and is a part of the UMA Python Client. When an `UMA_service` is created, it does not maintain a connection to the UMA core: it only stores the necessary information (url, port, header,...) necessary for a REST call to successfully communicate with UMA core.

An `UMA_service` object is a required input to any UMA Python Client object, and we recommend that the service object already present in the global scope be used for this purpose.

Though it is possible to create `UMA_service` objects on demand, the resulting temporary objects are only available within their parent scope, and require significant memory resources for their storage. In addition, multiple active `UMA_service` objects may interfere with each other during logging operations, which might pose a significant challenge when we graduate the UMA Python Client to a multi-threaded setting.

`UMA_service` is initialized with http scheme, host, and port designations. By default, it uses http, localhost and port 8000.

In addition to the `UMA_service` object, there are currently four python client classes that are used for REST calls:

Service_World: is initialized with an `UMA_service` instance, and is only responsible for constructing Agent objects;

Service_Agent: is initialized with an `UMA_service` instance and an agent id. It is responsible for BUA-level maintenance operations (e.g. adding snapshots, managing snapshot parameters, making a decision based on the internal state of member snapshots).

Service_Snapshot: is initialized with an `UMA_service` instance, an agent id (MID) and a snapshot id.³

The responsibilities of `Service_Snapshot` are quite diverse:

- (1) Adding existing sensors (by MIDs) to the snapshot;
- (2) Initializing snapshots (using the currently assigned sensors for initial sensors);
- (3) execute the up/ups/downs/propagation/abduction/npdir operations;
- (4) execute enrichment/pruning operations;
- (5) compute, store and report internal state vectors (current/prediction/target/negligible...);
- (6) set/update snapshot data in the UMA Core (q, threshold etc.).

Service_Sensor: is initiated with an `UMA_service` instance, an agent id (MID), snapshot id (same as above) and a sensor ID. It is responsible for obtaining the delay mask (if any) of the current sensor in the UMA Core.

To add a new REST functionality, define the corresponding UMA Python Client object (see below), and call the functions. Try to ensure all calls to the Python Client remain confined to `som2`.

³Either "plus" or "minus" in the current implementation.

2.7.2. *UMA Python Client Development.* UMA Python Client is the bridge between the UMA simulation engine and the UMA core. To add a new functionality in UMA Python Client, it is necessary to:

- (1) Find/construct the REST endpoint in UMA Core;
- (2) Design the input/output from the simulation engine.

All the functionalities currently implemented follow the following design format. First, the input to the new function, to be sent to the UMA core, is defined; as a rule, it is important that all input/output be shaped in terms of the Python base classes (e.g. lists, tuples).⁴ Next, a CRUD REST request to the UMA Core (out of the options provided by the `UMA_service` instance, see Section 2.5) is made. Finally, we always check the return message within the function for error handling. If everything is OK, return the data to UMA Python Simulation.

2.7.3. *UMA Core Service Development.* UMA Core Service development has five parts (see Section 2.2 for a detailed review).

UMAUtil: This is the basic module of the UMA architecture. Any functionality introduced into this module must ensure its independence from 3rd party libraries. For cross platform functionalities (file creation, time function etc), we recommend the use of macros to distinguish between platforms, when necessary. All new functionalities in `UMAUtil` should come equipped with thorough tests: please add new test cases in `UMAUtilTest`.

UMAKernel: This is the module containing all kernel code running on CUDA. it should only depend on `cuda` and `UMAUtil`. When adding a new source file to this package, be sure to change the build script on Linux, so that the device code is built and linked correctly.

UMACore: This module defines all UMA objects. When making changes to existing functionalities or adding new functionalities, ensure those changes are consistent with the content of the corresponding objects in the UMA Python simulation engine.

UMARest: The module providing the UMA Core with REST functionality. IT encapsulates the library `cpprestsdk` within itself. Therefore, when making alterations to this module, one must ensure all inputs/outputs of interface functions use only types provided in `std`. This is important because it can efficiently reduce 3rd party library dependencies on other modules, saving memory and, most critically, computation time.

UMA: The executable part of UMA Core Service. When a new REST handler or new REST mapping is needed, just add a new handler class or new functions to handle the new REST URL.

REFERENCES

- [1] D.P. Guralnik and D.E. Koditschek. The uma-som platform for general learning (technical report). *technical report*, 2017. https://github.com/kotmasha/kodlab-uma/AFRL_tech_report_2017.pdf.

⁴Generically, the UMA core should only be sent floating point values (such as the motivational signal or simulation parameters) and Boolean vectors/matrices (sensory signals).